

```
library(caret)
library(gbm)
library(e1071)
```

```
data(scats)
```

```
str(scats)
```

Output:

```
'data.frame': 110 obs. of 19 variables:
 $ Species : Factor w/ 3 levels "bobcat","coyote",...: 2 2 1 2 2 2 1 1 1 1 ...
 $ Month : Factor w/ 9 levels "April","August",...: 4 4 4 4 4 4 4 4 4 4 ...
 $ Year : int 2012 2012 2012 2012 2012 2012 2012 2012 2012 2012 ...
 $ Site : Factor w/ 2 levels "ANNU","YOLA": 2 2 2 2 2 2 1 1 1 1 ...
 $ Location : Factor w/ 3 levels "edge","middle",...: 1 1 2 2 1 1 3 3 3 2 ...
 $ Age : int 5 3 3 5 5 5 1 3 5 5 ...
 $ Number : int 2 2 2 2 4 3 5 7 2 1 ...
 $ Length : num 9.5 14.9 8.5 8.9 6.5 5.5 11 20.5 ...
 $ Diameter : num 25.7 25.4 18.8 18.1 20.7 21.2 15.7 21.9 17.5 18 ...
 $ Taper : num 41.9 37.1 16.5 24.7 20.1 28.5 8.2 19.3 29.1 21.4 ...
 $ TI : num 1.63 1.46 0.88 1.36 0.97 1.34 0.52 0.88 1.66 1.19 ...
 $ Mass : num 15.9 17.6 8.4 7.4 25.4 ...
 $ d13C : num -26.9 -29.6 -28.7 -20.1 -23.2 ...
 $ d15N : num 6.94 9.87 8.52 5.79 7.01 8.28 4.2 3.89 7.34 6.06 ...
 $ CN : num 8.5 11.3 8.1 11.5 10.6 9.5 4.5 5.6 5.8 7.7 ...
 $ ropey : int 0 0 1 1 0 1 1 0 0 1 ...
 $ segmented: int 0 0 1 0 1 0 1 1 1 1 ...
 $ flat : int 0 0 0 0 0 0 0 0 0 0 ...
 $ scrape : int 0 0 1 0 0 0 1 0 0 0 ...
 >
```

```
> # 1. Set the Species column as the target/outcome and convert it to numeric.
```

```
> outcomeName<-'Species'
```

```
> #Converting outcome variable to numeric
```

```
> scat$Species<-ifelse(scat$Species=="bobcat",0,ifelse(scat$Species=="coyote", 1, 2))
```

```
> # After converting outcome variable to numeric
```

```
> str(scat)
```

```
'data.frame': 110 obs. of 19 variables:
 $ Species : num 1 1 0 1 1 1 0 0 0 0 ...
 $ Month : Factor w/ 9 levels "April","August",...: 4 4 4 4 4 4 4 4 4 4 ...
 $ Year : int 2012 2012 2012 2012 2012 2012 2012 2012 2012 2012 ...
 $ Site : Factor w/ 2 levels "ANNU","YOLA": 2 2 2 2 2 2 1 1 1 1 ...
 $ Location : Factor w/ 3 levels "edge","middle",...: 1 1 2 2 1 1 3 3 3 2 ...
 $ Age : int 5 3 3 5 5 5 1 3 5 5 ...
```

```

$ Number : int 2 2 2 2 4 3 5 7 2 1 ...
$ Length : num 9.5 14 9 8.5 8 9 6 5.5 11 20.5 ...
$ Diameter : num 25.7 25.4 18.8 18.1 20.7 21.2 15.7 21.9 17.5 18 ...
$ Taper : num 41.9 37.1 16.5 24.7 20.1 28.5 8.2 19.3 29.1 21.4 ...
$ TI : num 1.63 1.46 0.88 1.36 0.97 1.34 0.52 0.88 1.66 1.19 ...
$ Mass : num 15.9 17.6 8.4 7.4 25.4 ...
$ d13C : num -26.9 -29.6 -28.7 -20.1 -23.2 ...
$ d15N : num 6.94 9.87 8.52 5.79 7.01 8.28 4.2 3.89 7.34 6.06 ...
$ CN : num 8.5 11.3 8.1 11.5 10.6 9 5.4 5.6 5.8 7.7 ...
$ ropey : int 0 0 1 1 0 1 1 0 0 1 ...
$ segmented: int 0 0 1 0 1 0 1 1 1 1 ...
$ flat : int 0 0 0 0 0 0 0 0 0 0 ...
$ scrape : int 0 0 1 0 0 0 1 0 0 0 ...
>

```

> # 2. Remove the Month, Year, Site, Location features.

```
> scat[,c('Month', 'Year', 'Site', 'Location')] <- list(NULL)
```

```
> str(scat)
```

```
'data.frame': 110 obs. of 15 variables:
```

```

$ Species : num 1 1 0 1 1 1 0 0 0 0 ...
$ Age : int 5 3 3 5 5 5 1 3 5 5 ...
$ Number : int 2 2 2 2 4 3 5 7 2 1 ...
$ Length : num 9.5 14 9 8.5 8 9 6 5.5 11 20.5 ...
$ Diameter : num 25.7 25.4 18.8 18.1 20.7 21.2 15.7 21.9 17.5 18 ...
$ Taper : num 41.9 37.1 16.5 24.7 20.1 28.5 8.2 19.3 29.1 21.4 ...
$ TI : num 1.63 1.46 0.88 1.36 0.97 1.34 0.52 0.88 1.66 1.19 ...
$ Mass : num 15.9 17.6 8.4 7.4 25.4 ...
$ d13C : num -26.9 -29.6 -28.7 -20.1 -23.2 ...
$ d15N : num 6.94 9.87 8.52 5.79 7.01 8.28 4.2 3.89 7.34 6.06 ...
$ CN : num 8.5 11.3 8.1 11.5 10.6 9 5.4 5.6 5.8 7.7 ...
$ ropey : int 0 0 1 1 0 1 1 0 0 1 ...
$ segmented: int 0 0 1 0 1 0 1 1 1 1 ...
$ flat : int 0 0 0 0 0 0 0 0 0 0 ...
$ scrape : int 0 0 1 0 0 0 1 0 0 0 ...
>

```

> # 3. Check if any values are null.

> # If there are, impute missing values using KNN.

```
> sum(is.na(scat))
```

```
[1] 47
```

```

> preProcValues <- preProcess(scat[,c('Age', 'Number', 'Length',
'Diameter', 'Taper', 'TI', 'Mass', 'd13C', 'd15N', 'CN', 'ropey', 'segmented', 'flat', 'scrape')], method =
c("knnImpute", "center", "scale"))
> library('RANN')

```

```
> scat_processed <- predict(preProcValues, scat)
> sum(is.na(scat_processed))
[1] 0
>
```

4. Converting every categorical variable to numerical (if needed).

```
> str(scat_processed) # after processing
'data.frame': 110 obs. of 15 variables:
 $ Species : num 1 1 0 1 1 1 0 0 0 0 ...
 $ Age : num 1.207 -0.252 -0.252 1.207 1.207 ...
 $ Number : num -0.433 -0.433 -0.433 -0.433 0.968 ...
 $ Length : num 0.0587 1.3679 -0.0867 -0.2322 -0.3777 ...
 $ Diameter : num 1.8396 1.7623 0.0622 -0.1181 0.5516 ...
 $ Taper : num 0.961 0.642 -0.726 -0.182 -0.487 ...
 $ TI : num 0.0283 -0.1406 -0.7171 -0.24 -0.6277 ...
 $ Mass : num 0.388 0.583 -0.458 -0.571 1.469 ...
 $ d13C : num 0.00468 -1.26856 -0.85947 3.12113 1.66403 ...
 $ d15N : num -0.165 0.807 0.359 -0.546 -0.141 ...
 $ CN : num 0.0276 0.7922 -0.0816 0.8468 0.6011 ...
 $ ropey : num -1.131 -1.131 0.876 0.876 -1.131 ...
 $ segmented: num -1.131 -1.131 0.876 -1.131 0.876 ...
 $ flat : num -0.239 -0.239 -0.239 -0.239 -0.239 ...
 $ scrape : num -0.217 -0.217 4.562 -0.217 -0.217 ...
>
```

Answer - After removing columns in Question 2 and converting outcome column
to numeric, there are no more categorical variables.
There is no need for any conversion at this point.

5. With a seed of 100, 75% training, 25% testing.
Build the following models: randomforest, neural net, naive bayes and GBM.
a. For these models display a) model summarization and
b) plot variable of importance, for the predictions (use the prediction set) display
c) confusion matrix (60 points)

```
> #Converting the dependent variable back to categorical
> scat_processed$Species<-as.factor(scat_processed$Species)
> str(scat_processed)
'data.frame': 110 obs. of 15 variables:
 $ Species : Factor w/ 3 levels "0","1","2": 2 2 1 2 2 2 1 1 1 1 ...
 $ Age : num 1.207 -0.252 -0.252 1.207 1.207 ...
 $ Number : num -0.433 -0.433 -0.433 -0.433 0.968 ...
```

```

$ Length : num 0.0587 1.3679 -0.0867 -0.2322 -0.3777 ...
$ Diameter : num 1.8396 1.7623 0.0622 -0.1181 0.5516 ...
$ Taper : num 0.961 0.642 -0.726 -0.182 -0.487 ...
$ TI : num 0.0283 -0.1406 -0.7171 -0.24 -0.6277 ...
$ Mass : num 0.388 0.583 -0.458 -0.571 1.469 ...
$ d13C : num 0.00468 -1.26856 -0.85947 3.12113 1.66403 ...
$ d15N : num -0.165 0.807 0.359 -0.546 -0.141 ...
$ CN : num 0.0276 0.7922 -0.0816 0.8468 0.6011 ...
$ ropey : num -1.131 -1.131 0.876 0.876 -1.131 ...
$ segmented: num -1.131 -1.131 0.876 -1.131 0.876 ...
$ flat : num -0.239 -0.239 -0.239 -0.239 -0.239 ...
$ scrape : num -0.217 -0.217 4.562 -0.217 -0.217 ...
>

```

```

> #Splitting training set into two parts based on outcome: 75% and 25%
> set.seed(100)
> index <- createDataPartition(scat_processed$Species, p=0.75, list=FALSE)
> trainSet <- scat_processed[ index,]
> testSet <- scat_processed[-index,]
> #Checking the structure of trainSet
> str(trainSet)
'data.frame': 83 obs. of 15 variables:
 $ Species : Factor w/ 3 levels "0","1","2": 2 1 2 2 2 1 1 3 3 3 ...
 $ Age : num 1.207 -0.252 1.207 1.207 1.207 ...
 $ Number : num -0.433 -0.433 -0.433 0.968 0.268 ...
 $ Length : num 0.0587 -0.0867 -0.2322 -0.3777 -0.0867 ...
 $ Diameter : num 1.8396 0.0622 -0.1181 0.5516 0.6804 ...
 $ Taper : num 0.9609 -0.7262 -0.1816 -0.4871 0.0709 ...
 $ TI : num 0.0283 -0.7171 -0.24 -0.6277 -0.2599 ...
 $ Mass : num 0.388 -0.458 -0.571 1.469 0.19 ...
 $ d13C : num 0.00468 -0.85947 3.12113 1.66403 -0.98357 ...
 $ d15N : num -0.165 0.359 -0.546 -0.141 0.28 ...
 $ CN : num 0.0276 -0.0816 0.8468 0.6011 0.1642 ...
 $ ropey : num -1.131 0.876 0.876 -1.131 0.876 ...
 $ segmented: num -1.131 0.876 -1.131 0.876 -1.131 ...
 $ flat : num -0.239 -0.239 -0.239 -0.239 -0.239 ...
 $ scrape : num -0.217 4.562 -0.217 -0.217 -0.217 ...
> # all variables as predictors
> predictors<-c("Age", "Number", "Length", "Diameter", "Taper", "TI", "Mass", "d13C",
+ "d15N", "CN", "ropey", "segmented", "flat", "scrape")
>

> # ##### randomforest #####
> model_rf<-train(trainSet[,predictors],trainSet[,outcomeName],method='rf', importance=T)

```

```
> # summarizing the model
> print(model_rf)
Random Forest
```

83 samples
14 predictors
3 classes: '0', '1', '2'


No pre-processing
Resampling: Bootstrapped (25 reps)
Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...
Resampling results across tuning parameters:

mtry	Accuracy	Kappa
2	0.6173904	0.3456397
8	0.6479826	0.4143276
14	0.6449064	0.4170319

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was mtry = 8.

```
>
```

```
> # Visualizing the models
> plot(model_rf)
>
```



```
> varImp(object=model_rf)
rf variable importance
```

variables are sorted by maximum importance across the classes

0	1	2
---	---	---

```


CN      100.000 15.788 84.801
Mass    34.267 25.760 66.057
d13C    64.580 60.099 37.362
d15N    44.807 51.268 34.111
segmented 40.202 0.000 48.987
Diameter 11.442 29.001 35.856
flat     24.498  5.886 35.177
Number   15.000 10.028 20.908
ropey    14.312 13.604 18.244
Taper    17.905 13.834 14.042
Length    1.228 17.413  9.645
TI        16.384  7.761 11.851
scrape   11.487 15.256 15.256
Age       13.491 12.742 13.537
>

```

```

> #Plotting Variable importance for Random Forest
> plot(varImp(object=model_rf),main="Random Forest - Variable Importance")
>

```



```

> #Predictions
> predictions_rf<-predict.train(object=model_rf,testSet[,predictors],type="raw")
> table(predictions_rf)
predictions_rf
 0  1  2
18  5  4
> #Confusion Matrix and Statistics
> confusionMatrix(predictions_rf,testSet[,outcomeName])
Confusion Matrix and Statistics

```

Reference

Prediction	0	1	2
0	14	2	2
1	0	5	0
2	0	0	4

Overall Statistics

Accuracy : 0.8519
 95% CI : (0.6627, 0.9581)
 No Information Rate : 0.5185
 P-Value [Acc > NIR] : 0.0003126

Kappa : 0.7416

Mcnemar's Test P-Value : NA

Statistics by Class:

	Class: 0	Class: 1	Class: 2
Sensitivity	1.0000	0.7143	0.6667
Specificity	0.6923	1.0000	1.0000
Pos Pred Value	0.7778	1.0000	1.0000
Neg Pred Value	1.0000	0.9091	0.9130
Prevalence	0.5185	0.2593	0.2222
Detection Rate	0.5185	0.1852	0.1481
Detection Prevalence	0.6667	0.1852	0.1481
Balanced Accuracy	0.8462	0.8571	0.8333

>

```
> # ##### Neural Net #####
> model_nnet<-train(trainSet[,predictors],trainSet[,outcomeName],method='nnet',
importance=T)
> # summarizing the model
> print(model_nnet)
Neural Network
```


83 samples
 14 predictors
 3 classes: '0', '1', '2'

No pre-processing
 Resampling: Bootstrapped (25 reps)
 Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...
 Resampling results across tuning parameters:

	size	decay	Accuracy	Kappa
1	0e+00	0.5506913	0.2859758	
1	1e-04	0.5285263	0.2388859	
1	1e-01	0.5720865	0.3012730	
3	0e+00	0.6198346	0.3975515	
3	1e-04	0.6497167	0.4392089	
3	1e-01	0.6850518	0.4921780	
5	0e+00	0.6559566	0.4472659	
5	1e-04	0.6495575	0.4435142	
5	1e-01	0.6943589	0.5100720	

Accuracy was used to select the optimal model using the largest value.
The final values used for the model were size = 5 and decay = 0.1.

```
> # Visualizing the models
> plot(model_nnet)
```



```
> #Variable Importance
> nnet_varImp<- varImp(object=model_nnet)
> #Plotting Variable importance for Neural Net
> plot(varImp(object=model_nnet),main="Neural Net - Variable Importance")
Error in rep.int(factor(names(x), unique(names(x))), lengths(x)) :
  invalid 'times' value
> #Predictions
> predictions_nnet<-predict.train(object=model_nnet,testSet[,predictors],type="raw")
> table(predictions_nnet)
predictions_nnet
 0  1  2
15  7  5
> #Confusion Matrix and Statistics
> confusionMatrix(predictions_nnet,testSet[,outcomeName])
```


Confusion Matrix and Statistics

Reference
Prediction 0 1 2
0 13 1 1
1 1 5 1
2 0 1 4

Overall Statistics

Accuracy : 0.8148
95% CI : (0.6192, 0.937)
No Information Rate : 0.5185
P-Value [Acc > NIR] : 0.001421

Kappa : 0.6932

McNemar's Test P-Value : 0.801252

Statistics by Class:

	Class: 0	Class: 1	Class: 2
Sensitivity	0.9286	0.7143	0.6667
Specificity	0.8462	0.9000	0.9524
Pos Pred Value	0.8667	0.7143	0.8000
Neg Pred Value	0.9167	0.9000	0.9091
Prevalence	0.5185	0.2593	0.2222
Detection Rate	0.4815	0.1852	0.1481
Detection Prevalence	0.5556	0.2593	0.1852
Balanced Accuracy	0.8874	0.8071	0.8095

>

```
> # ##### Naive Bayes #####  
> model_nb<-train(trainSet[,predictors],trainSet[,outcomeName],method='naive_bayes',  
importance=T)  
There were 50 or more warnings (use warnings() to see the first 50)  
> # summarizing the model  
> print(model_nb)  
Naive Bayes  
  
83 samples  
14 predictors  
3 classes: '0', '1', '2'
```

No pre-processing

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...

Resampling results across tuning parameters:

usekernel	Accuracy	Kappa
FALSE	0.5971475	0.3907572
TRUE	0.6461044	0.4180012

Tuning parameter 'laplace' was held constant at a value of 0

Tuning parameter 'adjust'

was held constant at a value of 1


Accuracy was used to select the optimal model using the largest value.

The final values used for the model were laplace = 0, usekernel = TRUE and adjust = 1.

> # Visualizing the models

> plot(model_nb)

>



> #Variable Importance

> varImp(object=model_nb)

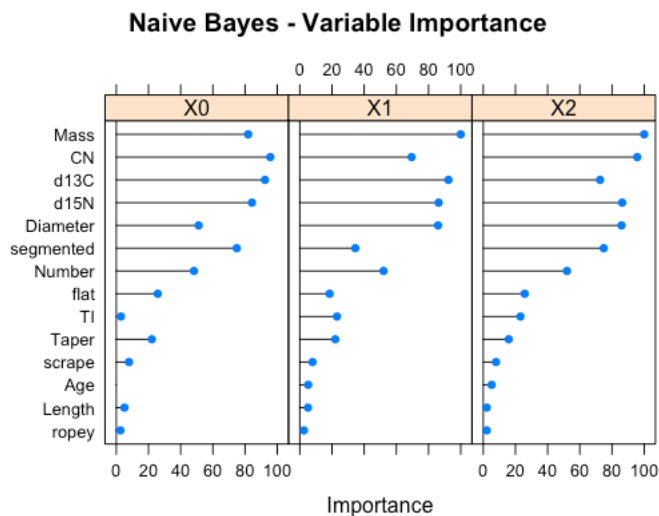
ROC curve variable importance

variables are sorted by maximum importance across the classes

	X0	X1	X2
Mass	81.911	100.000	100.000
CN	95.670	69.479	95.670
d13C	92.359	92.359	72.587
d15N	84.284	86.294	86.294
Diameter	51.142	85.913	85.913
segmented	74.845	34.487	74.845
Number	48.256	52.028	52.028

```
flat 25.757 18.523 25.757
TI 2.860 23.092 23.092
Taper 22.038 22.038 15.858
scrape 7.907 7.907 7.907
Age 0.000 5.197 5.197
Length 5.047 5.047 2.152
ropey 2.523 2.523 2.143
```

```
> #Plotting Variable importance for Naive Bayes
> plot(varImp(object=model_nb),main="Naive Bayes - Variable Importance")
>
```



```
> #Predictions
> predictions_nb<-predict.train(object=model_nb,testSet[,predictors],type="raw")
> table(predictions_nb)
predictions_nb
0 1 2
18 5 4
> #Confusion Matrix and Statistics
> confusionMatrix(predictions_nb,testSet[,outcomeName])
Confusion Matrix and Statistics
```

```
Reference
Prediction 0 1 2
0 14 2 2
1 0 5 0
2 0 0 4
```

Overall Statistics

```
Accuracy : 0.8519
95% CI : (0.6627, 0.9581)
```

No Information Rate : 0.5185
P-Value [Acc > NIR] : 0.0003126

Kappa : 0.7416

McNemar's Test P-Value : NA

Statistics by Class:

	Class: 0	Class: 1	Class: 2
Sensitivity	1.0000	0.7143	0.6667
Specificity	0.6923	1.0000	1.0000
Pos Pred Value	0.7778	1.0000	1.0000
Neg Pred Value	1.0000	0.9091	0.9130
Prevalence	0.5185	0.2593	0.2222
Detection Rate	0.5185	0.1852	0.1481
Detection Prevalence	0.6667	0.1852	0.1481
Balanced Accuracy	0.8462	0.8571	0.8333

>

> # ##### GBM #####

> model_gbm1<-train(trainSet[,predictors],trainSet[,outcomeName],method='gbm')

> # summarizing the model

> print(model_gbm1)

Stochastic Gradient Boosting

83 samples

14 predictors

3 classes: '0', '1', '2'

No pre-processing

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...

Resampling results across tuning parameters:

interaction.depth	n.trees	Accuracy	Kappa
1	50	0.6147733	0.3651472
1	100	0.6155064	0.3710730
1	150	0.6088021	0.3618473
2	50	0.5941397	0.3363546
2	100	0.6021387	0.3509252
2	150	0.5993344	0.3452358
3	50	0.5977627	0.3374063

3	100	0.6070331	0.3523725
3	150	0.6008749	0.3491138

Tuning parameter 'shrinkage' was held constant at a value of 0.1

Tuning

parameter 'n.minobsinnode' was held constant at a value of 10


Accuracy was used to select the optimal model using the largest value.

The final values used for the model were n.trees = 100, interaction.depth = 1, shrinkage = 0.1 and n.minobsinnode = 10.

> # Visualizing the models

> plot(model_gbm1)

>



> #Variable Importance

> varImp(object=model_gbm1)


gbm variable importance

Overall	
CN	100.000
d13C	74.587
Mass	63.478
Diameter	54.631
d15N	41.043
TI	27.289
Length	24.577
Taper	22.064
Number	14.108
segmented	10.024
Age	5.829
ropey	5.020
scrape	0.000

flat 0.000

> #Plotting Variable importance for GBM

> plot(varImp(object=model_gbm1),main="GBM - Variable Importance")



> #Predictions

> predictions_gbm<-predict.train(object=model_gbm1,testSet[,predictors],type="raw")

> table(predictions_gbm)

predictions_gbm

0 1 2

17 5 5

> #Confusion Matrix and Statistics

> confusionMatrix(predictions_gbm,testSet[,outcomeName])

Confusion Matrix and Statistics

Reference

Prediction 0 1 2

0 14 1 2

1 0 5 0

2 0 1 4

Overall Statistics

Accuracy : 0.8519

95% CI : (0.6627, 0.9581)

No Information Rate : 0.5185

P-Value [Acc > NIR] : 0.0003126

Kappa : 0.7465

Mcnemar's Test P-Value : 0.2614641

Statistics by Class:

```
          Class: 0 Class: 1 Class: 2
Sensitivity      1.0000  0.7143  0.6667
Specificity      0.7692  1.0000  0.9524
Pos Pred Value   0.8235  1.0000  0.8000
Neg Pred Value   1.0000  0.9091  0.9091
Prevalence       0.5185  0.2593  0.2222
Detection Rate   0.5185  0.1852  0.1481
Detection Prevalence 0.6296  0.1852  0.1852
Balanced Accuracy 0.8846  0.8571  0.8095
>
```

```
# 6. For the BEST performing models of each (randomforest, neural net, naive bayes and gbm)
# create and display a data frame that has the following columns:
# ExperimentName, accuracy, kappa.
# Sort the data frame by accuracy.
```

```
> experimentName<-c("Random Forest", "Neural Net", "Naive Bayes", "GBM")
> accuracyDetails<-c(max(model_rf$results$Accuracy), max(model_nnet$results$Accuracy),
+                    max(model_nb$results$Accuracy), max(model_gbm1$results$Accuracy))
> kappaDetails<-c(max(model_rf$results$Kappa), max(model_nnet$results$Kappa),
+                 max(model_nb$results$Kappa), max(model_gbm1$results$Kappa))
> bestModelDf<-data.frame(ExperimentName=experimentName, Accuracy=accuracyDetails,
Kappa=kappaDetails)
> print(bestModelDf[order(-bestModelDf$Accuracy),])
  ExperimentName Accuracy  Kappa
2   Neural Net 0.6943589 0.5100720
1 Random Forest 0.6479826 0.4170319
3   Naive Bayes 0.6461044 0.4180012
4         GBM 0.6155064 0.3710730
>
```

```
> # 7. Tune the GBM model using tune length = 20 and:
> # a) print the model summary and b) plot the models. (20 points)
> #using tune length
> fitControl <- trainControl(
+   method = "repeatedcv",
+   number = 5,
+   repeats = 5)
> #### Using tuneGrid ####
> modelLookup(model='gbm')
  model      parameter      label forReg forClass probModel
```

```

1 gbm      n.trees # Boosting Iterations TRUE  TRUE  TRUE
2 gbm interaction.depth      Max Tree Depth TRUE  TRUE  TRUE
3 gbm      shrinkage      Shrinkage TRUE  TRUE  TRUE
4 gbm  n.minobsinnode Min. Terminal Node Size TRUE  TRUE  TRUE
> model_gbm2<-
train(trainSet[,predictors],trainSet[,outcomeName],method='gbm',trControl=fitControl,tuneLength=20)
> # a) print the model summary
> print(model_gbm2)
Stochastic Gradient Boosting

```

```

83 samples
14 predictors
3 classes: '0', '1', '2'

```

```

No pre-processing
Resampling: Cross-Validated (5 fold, repeated 5 times)
Summary of sample sizes: 66, 66, 66, 67, 67, 67, ...
Resampling results across tuning parameters:


```

	interaction.depth	n.trees	Accuracy	Kappa
1	50	0.6576340	0.4311423	
1	100	0.6417516	0.4060450	
1	150	0.6124150	0.3621097	
1	200	0.6171209	0.3680670	
1	250	0.6100621	0.3549479	
1	300	0.6077255	0.3583588	
1	350	0.5980033	0.3379037	

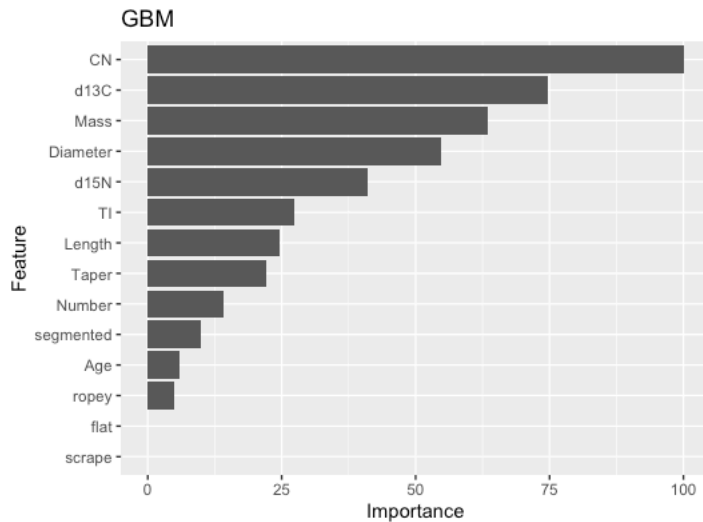
```

....
> #b) plot the models.
> plot(model_gbm2)
>

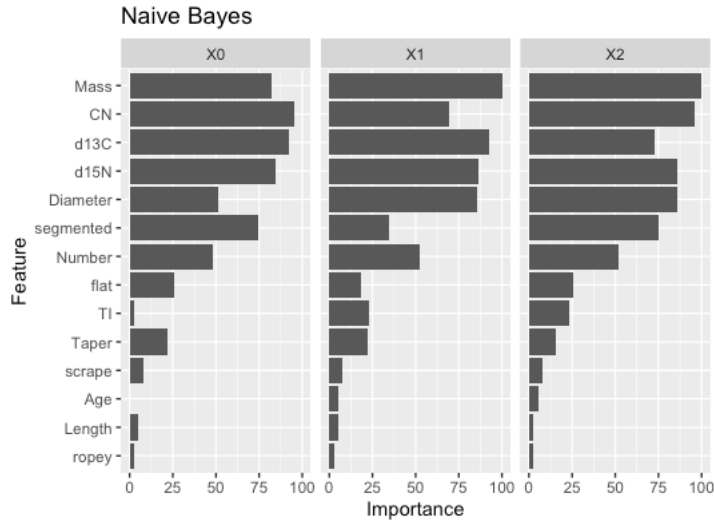
```

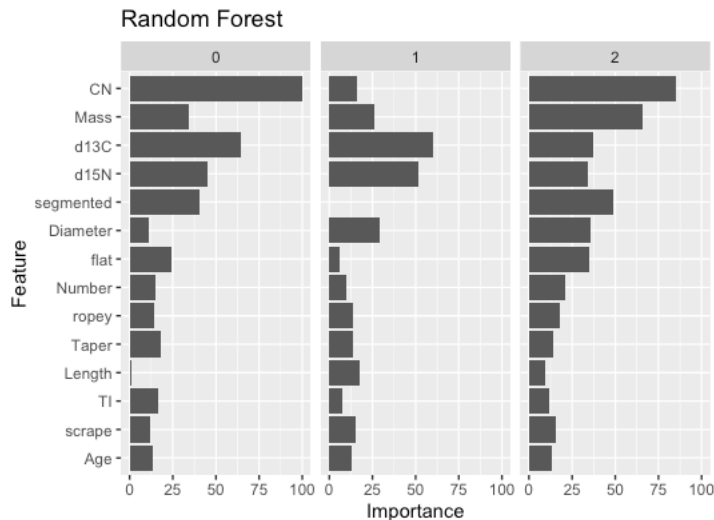
```
> # 8. Using GGplot and gridExtra to plot all variable of
> # importance plots into one single plot. (10 points)
> library(ggplot2)
> library(gridExtra)
> plot_gbm1 <- ggplot(data=varImp(object=model_gbm1)) + ggtitle("GBM")
> print(plot_gbm1)
```



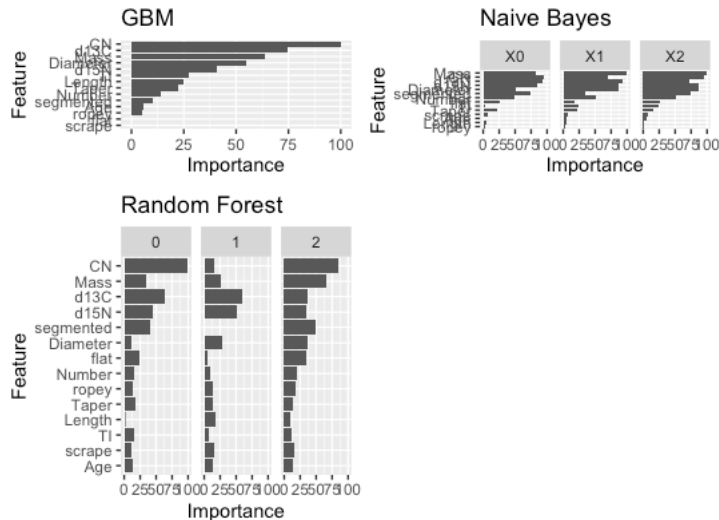
```
> plot_nb <- ggplot(data=varImp(object=model_nb)) + ggtitle("Naive Bayes")
> print(plot_nb)
```



```
> plot_nnet <- ggplot(data=varImp(object=model_nnet)) + ggtitle("Neural Network")
Error in rep.int(factor(names(x), unique(names(x))), lengths(x)) :
  invalid 'times' value
> print(plot_nnet)
Error in print(plot_nnet) : object 'plot_nnet' not found
> plot_rf <- ggplot(data=varImp(object=model_rf)) + ggtitle("Random Forest")
> print(plot_rf)
```



```
> grid.arrange(plot_gbm1, plot_nb, plot_rf, nrow = 2, ncol=2, heights = c(0.35, 0.65))
```



9. Which model performs the best? and why do you think this is the case?

Can we accurately predict species on this dataset? (10 points)

Answer -

As per accuracy & Kappa values of the models, Neural Network model performs the best.

ExperimentName Accuracy Kappa

2 Neural Net 0.6943589 0.5100720

1 Random Forest 0.6479826 0.4170319

3 Naive Bayes 0.6461044 0.4180012

4 GBM 0.6155064 0.3710730

Neural Network model uses memory to store previous read values and

does classification based on correlation between the read values.

Whereas, other models like Naive Bayes considers attributes are independent of each other.

Species from the data set can be predicted with less than 70 % accuracy

i.e. 69.4% using Neural Networks model, 64.8% using Random Forest model can be accurately predicted.

10. Graduate Student questions:

a. Using feature selection with rfe in caret and the repeatedcv method: Find the top 3

predictors and build the same models as in 6 and 8 with the same parameters. (20 points)

> #Feature selection using rfe in caret

> control <- rfeControl(functions = rfFuncs,

+ method = "repeatedcv",

+ repeats = 3,

+ verbose = FALSE)

> predictors<-names(trainSet)[!names(trainSet) %in% outcomeName]

> Species_Pred <- rfe(trainSet[,predictors], trainSet[,outcomeName],rfeControl = control)

> Species_Pred

Recursive feature selection

Outer resampling method: Cross-Validated (10 fold, repeated 3 times)

Resampling performance over subset size:

Variables	Accuracy	Kappa	AccuracySD	KappaSD	Selected
4	0.7012	0.4856	0.1523	0.2653	*
8	0.6927	0.4629	0.1643	0.2865	
14	0.6870	0.4461	0.1518	0.2794	

The top 4 variables (out of 4):

CN, d13C, d15N, Mass

```
> #Taking only the top 3 predictors
> predictorsTop3<-c("CN", "d13C", "d15N")
> # ##### randomforest with top-3 selected features #####
> model_rf_new<-train(trainSet[,predictorsTop3],trainSet[,outcomeName],method='rf',
importance=T)
note: only 2 unique complexity parameters in default grid. Truncating the grid to 2 .
```

```
> # summarizing the model
> print(model_rf_new)
Random Forest
```

83 samples
3 predictor
3 classes: '0', '1', '2'

No pre-processing

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...


Resampling results across tuning parameters:

mtry	Accuracy	Kappa
2	0.6797157	0.4496038
3	0.6656016	0.4308129

Accuracy was used to select the optimal model using the largest value.

The final value used for the model was mtry = 2.

```
> # Visualizing the models
> plot(model_rf_new)
```




```
> #Variable Importance
> varImp(object=model_rf_new)
rf variable importance
```

variables are sorted by maximum importance across the classes

```
0 1 2
CN 100.00 0.00 55.615
d13C 42.83 24.03 9.208
d15N 17.08 28.46 2.506
```

```
> #Plotting Variable importance for Random Forest
> plot(varImp(object=model_rf_new),main="Random Forest - Variable Importance")
```



```
> #Predictions
> predictions_rf_new<-
predict.train(object=model_rf_new,testSet[,predictorsTop3],type="raw")
> table(predictions_rf_new)
predictions_rf_new
0 1 2
```

15 6 6

> #Confusion Matrix and Statistics

> confusionMatrix(predictions_rf_new,testSet[,outcomeName])

Confusion Matrix and Statistics

Reference

Prediction 0 1 2

0 11 1 3

1 1 5 0

2 2 1 3

Overall Statistics

Accuracy : 0.7037

95% CI : (0.4982, 0.8625)

No Information Rate : 0.5185

P-Value [Acc > NIR] : 0.04012

Kappa : 0.5102

Mcnemar's Test P-Value : 0.75300

Statistics by Class:

Class: 0 Class: 1 Class: 2

Sensitivity 0.7857 0.7143 0.5000

Specificity 0.6923 0.9500 0.8571

Pos Pred Value 0.7333 0.8333 0.5000

Neg Pred Value 0.7500 0.9048 0.8571

Prevalence 0.5185 0.2593 0.2222

Detection Rate 0.4074 0.1852 0.1111

Detection Prevalence 0.5556 0.2222 0.2222

Balanced Accuracy 0.7390 0.8321 0.6786

> # ##### Neural Net with top-3 selected features #####

> model_nnet_new<-train(trainSet[,predictorsTop3],trainSet[,outcomeName],method='nnet',
importance=T)

> # summarizing the model

> print(model_nnet_new)

Neural Network

83 samples

3 predictor

3 classes: '0', '1', '2'

No pre-processing

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...

Resampling results across tuning parameters:


	size	decay	Accuracy	Kappa
1	0e+00	0.6748811	0.4172198	
1	1e-04	0.6594939	0.3943782	
1	1e-01	0.6708776	0.4121009	
3	0e+00	0.6815514	0.4592315	
3	1e-04	0.6884762	0.4660593	
3	1e-01	0.7539512	0.5732017	
5	0e+00	0.6368239	0.3989197	
5	1e-04	0.6364624	0.3947993	
5	1e-01	0.7553847	0.5772432	

Accuracy was used to select the optimal model using the largest value.

The final values used for the model were size = 5 and decay = 0.1.

> # Visualizing the models

> plot(model_nnet_new)



> #Plotting Variable importance for Neural Net

> plot(varImp(object=model_nnet_new),main="Neural Net - Variable Importance")

Error in rep.int(factor(names(x), unique(names(x))), lengths(x)) :

invalid 'times' value

> #Predictions

> predictions_nnet_new<-

predict.train(object=model_nnet_new,testSet[,predictorsTop3],type="raw")

> table(predictions_nnet_new)

predictions_nnet_new

0 1 2

17 5 5

> #Confusion Matrix and Statistics

> confusionMatrix(predictions_nnet_new,testSet[,outcomeName])

Confusion Matrix and Statistics

Reference

Prediction 0 1 2

0 13 1 3

1 1 4 0

2 0 2 3

Overall Statistics

Accuracy : 0.7407

95% CI : (0.5372, 0.8889)

No Information Rate : 0.5185

P-Value [Acc > NIR] : 0.01571

Kappa : 0.5563

McNemar's Test P-Value : 0.17180

Statistics by Class:

Class: 0 Class: 1 Class: 2

Sensitivity 0.9286 0.5714 0.5000

Specificity 0.6923 0.9500 0.9048

Pos Pred Value 0.7647 0.8000 0.6000

Neg Pred Value 0.9000 0.8636 0.8636

Prevalence 0.5185 0.2593 0.2222

Detection Rate 0.4815 0.1481 0.1111

Detection Prevalence 0.6296 0.1852 0.1852

Balanced Accuracy 0.8104 0.7607 0.7024

>> # ##### Naive Bayes with top-3 selected features #####

> model_nb_new<-

train(trainSet[,predictorsTop3],trainSet[,outcomeName],method='naive_bayes', importance=T)

There were 50 or more warnings (use warnings() to see the first 50)

> # summarizing the model

> print(model_nb_new)

Naive Bayes

83 samples

3 predictor

3 classes: '0', '1', '2'

No pre-processing

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...

Resampling results across tuning parameters:

usekernel	Accuracy	Kappa
FALSE	0.7127563	0.4979608
TRUE	0.6867141	0.4609503

Tuning parameter 'laplace' was held constant at a value of 0

Tuning parameter 'adjust'


was held constant at a value of 1

Accuracy was used to select the optimal model using the largest value.

The final values used for the model were laplace = 0, usekernel = FALSE and adjust = 1.

> # Visualizing the models

> plot(model_nb_new)



> varImp(object=model_nb_new)


ROC curve variable importance

variables are sorted by maximum importance across the classes

	X0	X1	X2
CN	100.00	0.00	100.00
d13C	87.36	87.36	11.87
d15N	56.52	64.20	64.20

> #Plotting Variable importance for Naive Bayes

> plot(varImp(object=model_nb_new),main="Naive Bayes - Variable Importance")



```
> #Predictions
> predictions_nb_new<-
predict.train(object=model_nb_new,testSet[,predictorsTop3],type="raw")
> table(predictions_nb_new)
predictions_nb_new
 0  1  2
18  5  4
> #Confusion Matrix and Statistics
> confusionMatrix(predictions_nb_new,testSet[,outcomeName])
Confusion Matrix and Statistics
```

```
Reference
Prediction 0 1 2
0 14 1 3
1 0 5 0
2 0 1 3
```

Overall Statistics

```
Accuracy : 0.8148
95% CI : (0.6192, 0.937)
No Information Rate : 0.5185
P-Value [Acc > NIR] : 0.001421
```

```
Kappa : 0.677
```

```
Mcnemar's Test P-Value : 0.171797
```

Statistics by Class:

	Class: 0	Class: 1	Class: 2
Sensitivity	1.0000	0.7143	0.5000
Specificity	0.6923	1.0000	0.9524
Pos Pred Value	0.7778	1.0000	0.7500
Neg Pred Value	1.0000	0.9091	0.8696
Prevalence	0.5185	0.2593	0.2222
Detection Rate	0.5185	0.1852	0.1111
Detection Prevalence	0.6667	0.1852	0.1481
Balanced Accuracy	0.8462	0.8571	0.7262

> # ##### GBM with top-3 selected features #####

> model_gbm1_new<-train(trainSet[,predictorsTop3],trainSet[,outcomeName],method='gbm')

> # summarizing the model

> print(model_gbm1_new)

Stochastic Gradient Boosting

83 samples

3 predictor

3 classes: '0', '1', '2'

No pre-processing

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...

Resampling results across tuning parameters:

interaction.depth	n.trees	Accuracy	Kappa
1	50	0.6376478	0.4006285
1	100	0.6196917	0.3705796
1	150	0.6114344	0.3551965
2	50	0.6270669	0.3856325
2	100	0.6095393	0.3571104
2	150	0.6120484	0.3609020
3	50	0.6270374	0.3859090
3	100	0.6052552	0.3501068
3	150	0.6155551	0.3654858

Tuning parameter 'shrinkage' was held constant at a value of 0.1

Tuning


parameter 'n.minobsinnode' was held constant at a value of 10

Accuracy was used to select the optimal model using the largest value.

The final values used for the model were n.trees = 50, interaction.depth = 1, shrinkage = 0.1 and n.minobsinnode = 10.


> # Visualizing the models

> plot(model_gbm1_new)



```
> #Variable Importance
> varImp(object=model_gbm1_new)
gbm variable importance
```

```
Overall
d13C 100.00
CN 83.19
d15N 0.00
> #Plotting Variable importance for GBM
> plot(varImp(object=model_gbm1_new),main="GBM - Variable Importance")
GBM - Variable Importance
```



```
> #Predictions
> predictions_gbm_new<-
predict.train(object=model_gbm1_new,testSet[,predictorsTop3],type="raw")
> table(predictions_gbm_new)
predictions_gbm_new
0 1 2
17 7 3
```

```
> #Confusion Matrix and Statistics
> confusionMatrix(predictions_gbm_new,testSet[,outcomeName])
Confusion Matrix and Statistics
```

```

      Reference
Prediction 0 1 2
0 13 1 3
1 0 5 2
2 1 1 1
```

Overall Statistics

```

      Accuracy : 0.7037
      95% CI : (0.4982, 0.8625)
No Information Rate : 0.5185
P-Value [Acc > NIR] : 0.04012
```

```
      Kappa : 0.4906
```

```
McNemar's Test P-Value : 0.50617
```

Statistics by Class:

	Class: 0	Class: 1	Class: 2
Sensitivity	0.9286	0.7143	0.16667
Specificity	0.6923	0.9000	0.90476
Pos Pred Value	0.7647	0.7143	0.33333
Neg Pred Value	0.9000	0.9000	0.79167
Prevalence	0.5185	0.2593	0.22222
Detection Rate	0.4815	0.1852	0.03704
Detection Prevalence	0.6296	0.2593	0.11111
Balanced Accuracy	0.8104	0.8071	0.53571

```
# 10. b. Create a dataframe that compares the non-feature selected models ( the same as on 7)
# and add the best BEST performing models of each (randomforest, neural net, naive bayes and
gbm) and
```

```
# display the data frame that has the following columns: ExperimentName, accuracy, kappa.
```

```
# Sort the data frame by accuracy. (40 points)
```

```
> experimentName_new<-c("Random Forest New" , "Neural Net New", "Naive Bayes New",
"GBM New")
```

```
> accuracyDetails_new<-c(max(model_rf_new$results$Accuracy),
max(model_nnet_new$results$Accuracy),
+       max(model_nb_new$results$Accuracy),
max(model_gbm1_new$results$Accuracy))
```

```
> kappaDetails_new<-c(max(model_rf_new$results$Kappa),
max(model_nnet_new$results$Kappa),
+      max(model_nb_new$results$Kappa), max(model_gbm1_new$results$Kappa))
> bestModelDf_new<-data.frame(ExperimentName=c(experimentName,
experimentName_new),
+      Accuracy=c(accuracyDetails, accuracyDetails_new),
+      Kappa=c(kappaDetails, kappaDetails_new))
> print(bestModelDf_new[order(-bestModelDf_new$Accuracy),])
```

	ExperimentName	Accuracy	Kappa
6	Neural Net New	0.7553847	0.5772432
7	Naive Bayes New	0.7127563	0.4979608
2	Neural Net	0.6943589	0.5100720
5	Random Forest New	0.6797157	0.4496038
1	Random Forest	0.6479826	0.4170319
3	Naive Bayes	0.6461044	0.4180012
8	GBM New	0.6376478	0.4006285
4	GBM	0.6155064	0.3710730

NOTE::: "New" refers to the latest models with Top 3 predictors in
"Random Forest New" , "Neural Net New", "Naive Bayes New", "GBM New".

c. Which model performs the best? and why do you think this is the case?
Can we accurately predict species on this dataset? (10 points)

Answer -

Output of Best Model dataframe :::

#	ExperimentName	Accuracy	Kappa
# 6	Neural Net New	0.7522766	0.5693890
# 7	Naive Bayes New	0.7254481	0.5135670
# 2	Neural Net	0.6943589	0.5100720
# 5	Random Forest New	0.6861781	0.4607807
# 1	Random Forest	0.6479826	0.4170319
# 3	Naive Bayes	0.6461044	0.4180012
# 8	GBM New	0.6352902	0.3809075
# 4	GBM	0.6155064	0.3710730

From the Best model printed above with based on accuracy, it is evident that Neural Network
New

(with top 3 predictors) performed the best with an accuracy of 75.22% and Kappa value of
56.9%.

Followed by Naive Bayes(with top 3 predictors) and the old Neural network model (with non-
feature selected)

with accuracies of 72.54 & 69.43 % respectively.

With the new Neural Network model with top 3 selected features, prediction accuracy increased from 69.4 % to 75.22 %.

Prediction using this model is improved.

Yes we can predict Species from the dataset using Neural Network New model with 75.22% accuracy.